

Logistic Regression

Objective: "What sorts of people were more likely to survive the sinking of the Titanic?"

<https://www.kaggle.com/competitions/titanic/data>

```
[1]: import pandas as pd
import statsmodels.api as sm
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.metrics import roc_curve, roc_auc_score
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: t_df = pd.read_csv('titanic_data.csv', index_col='PassengerId')
```

```
[3]: print(t_df.shape)

# Count null values per column
# null_counts = t_df.isnull().sum()
# print(null_counts)

# Drop all rows containing at least one null value
# t_df = t_df.dropna()
# print(t_df.shape)

(891, 11)
```

```
[10]: #Data Cleanup

cols_to_drop = ['Name', 'Cabin', 'Ticket']
t_df.drop(columns=[col for col in cols_to_drop if col in t_df.columns], inplace=True)

t_df = t_df.dropna()

t_df['Sex'] = t_df['Sex'].replace({'male': 1, 'female': 0})
```

```
[11]: X = t_df.drop(columns=['Survived'])
y = t_df['Survived']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
[12]: X_train_sm = sm.add_constant(X_train)

X_train_sm = X_train_sm.dropna()
y_train_cleaned = y_train.loc[X_train_sm.index]

logmodel = sm.Logit(y_train_cleaned, X_train_sm).fit()

print(logmodel.summary())
```

Optimization terminated successfully.
Current function value: 0.413088
Iterations 7

Logit Regression Results

```
=====
Dep. Variable:    Survived    No. Observations:    498
Model:            Logit      Df Residuals:        490
Method:           MLE        Df Model:            7
Date:             Wed, 23 Apr 2025    Pseudo R-squ.:      0.3837
Time:             23:50:03    Log-Likelihood:     -205.72
converged:        True        LL-Null:            -333.82
Covariance Type:  nonrobust    LLR p-value:        1.324e-51
=====
```

	coef	std err	z	P> z	[0.025	0.975]
const	6.0165	0.787	7.643	0.000	4.474	7.559
Pclass	-1.5647	0.219	-7.129	0.000	-1.995	-1.135
Sex	-2.8457	0.284	-10.033	0.000	-3.402	-2.290
Age	-0.0410	0.010	-4.053	0.000	-0.061	-0.021
SibSp	-0.2944	0.155	-1.895	0.058	-0.599	0.010
Parch	0.0266	0.161	0.165	0.869	-0.289	0.342
Fare	-0.0008	0.004	-0.210	0.834	-0.009	0.007
Embarked	0.1498	0.257	0.582	0.561	-0.354	0.654

=====

```
[15]: X_test_sm = sm.add_constant(X_test)
X_test_sm = X_test_sm.dropna()
y_test_cleaned = y_test.loc[X_test_sm.index]

pred_probs = logmodel.predict(X_test_sm)

predictions = [1 if p > 0.5 else 0 for p in pred_probs]


from sklearn.metrics import accuracy_score, confusion_matrix
print("Accuracy:", accuracy_score(y_test_cleaned, pred_binary))
print("Confusion Matrix:\n", confusion_matrix(y_test_cleaned, pred_binary))
```

```
Accuracy: 0.780373831775701
Confusion Matrix:
[[105  17]
 [ 30  62]]
```

```
[16]: fpr, tpr, thresholds = roc_curve(y_test, predictions)
roc_auc = roc_auc_score(y_test, predictions)

plt.plot(fpr, tpr, label='ROC Curve (area = %0.3f)' % roc_auc)
plt.title('ROC Curve (area = %0.3f)' % roc_auc)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
```

```
[16]: Text(0, 0.5, 'True Positive Rate')
```



Task 1: Rearrange the code so that the cabin column is removed BEFORE the null values. How does this change the confusion matrix and overall accuracy?

```
[17]: df = pd.read_csv('titanic_data.csv', index_col='PassengerId')

[18]: df.drop(columns=['Cabin', 'Name', 'Ticket'], inplace=True)

[19]: df = df.dropna()

[21]: df['Sex'] = df['Sex'].replace({'male': 1, 'female': 0}).astype(int)
      df['Embarked'] = df['Embarked'].replace({'S': 0, 'C': 1, 'Q': 2}).astype(int)

[22]: X = df.drop(columns='Survived')
      y = df['Survived']

[23]: from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
[24]: import statsmodels.api as sm
      X_train_sm = sm.add_constant(X_train).dropna()
      y_train_sm = y_train.loc[X_train_sm.index]
      logmodel = sm.Logit(y_train_sm, X_train_sm).fit(dispatch=False)

[25]: X_test_sm = sm.add_constant(X_test).dropna()
      y_test_cleaned = y_test.loc[X_test_sm.index]
      pred_probs = logmodel.predict(X_test_sm)
      predictions = [1 if p > 0.5 else 0 for p in pred_probs]

[26]: from sklearn.metrics import accuracy_score, confusion_matrix
      print("Confusion Matrix:\n", confusion_matrix(y_test_cleaned, predictions))
      print("Accuracy: %.2f%%" % (accuracy_score(y_test_cleaned, predictions) * 100))
```

```
Confusion Matrix:
[[105  17]
 [ 30  62]]
Accuracy: 78.04%
```